

CS 580 Homework 7

Kevin Mao

May 1, 2015

Background

The task of reliable and robust facial recognition represents a number of challenging computational, engineering, and practical problems. *Eigenfaces for Recognition* by Matthew Turk and Alex Pentland proposes a novel solution to this problem by deriving common vectors between a set of training faces, named Eigenfaces. These Eigenfaces are used to recognize new faces by calculating the magnitude of distance between the new image's representative image vector and the image vectors of the Eigenfaces.

This assignment recreates Turk and Pentland's original facial recognition algorithm, trains it using a randomly selected number of facial images from a well known database, and evaluates recognition against a second set of randomly selected facial images.

Generating Eigenfaces

Let each image be a training set of M distinct N by N pixel resolution images. For the purposes of simplicity, we convert each of these images into greyscale intensity values. Accounting for color during Eigenface generation is an area of future enhancement. Each of these N by N images are represented as a one dimensional array of values of length N^2 .

Let the set of one dimensional arrays representing the greyscale training face images be represented as (1).

$$\tau = \tau_1, \tau_2, \dots, \tau_M \tag{1}$$

The average face image (2) can be defined as a one dimensional array of size N with the value at each index equal to the average of pixel value of the training faces at that index.

$$\psi = \frac{1}{M} \sum_{n=1}^M (\tau_n) \quad (2)$$

We can then calculate a set of nomralized pixel vectors (3) representing each training face image

$$\phi_i = \tau_i - \psi \quad (3)$$

The naive approach to generating Eigenfaces from these normalized vectors would be to generate the covariance matrix (4) and derive the resultant N^2 eigenvalue/eigenvector pairs.

$$C = \frac{1}{M} \sum_{n=1}^M (\phi_n \phi_n^T) \quad (4)$$

$$= AA^T \quad (5)$$

This approach, however, becomes intractable as the generated covariance matrix is of size N^2 by N^2 . For example, given a set of 256 by 256 greyscale training face images, where each pixel intensity value is represented as a 4-byte integer, the covariance matrix alone would require 16 gigabytes of memory. Turk and Pentland propose a more scalable solution based on the assumption that if $M \ll N^2$, only $M - 1$ out of N eigenvectors will be meaningful. Consider the eigenvalues μ_i and eigenvectors v_i of $A^T A$:

$$A^T A t_i = \mu_i t_i \quad (6)$$

Premultiplying both sides of the equation yields (6), which is equivalent to the the eigenvalue/eigenvector pairs of (4), where the eigenvector is equal to Av_i .

$$AA^T Av_i = \mu_i Av_i \quad (7)$$

The set of eigenvalue/eigenvector pairs (μ_i, Av_i) has a size of N^2 . However only $(M - 1)$ pairs have a nonzero eigenvalue. From these $(M - 1)$ we select the M' pairs containing the largest eigenvalues, where $M' < (M - 1)$ is some number of eigenfaces to use for identification/recognition. These form the set of eigenfaces u_1, u_2, \dots, u'_M .

Classifying a Face Image

Recognizing a Face Image

Any input image, τ , can be transformed into a set of weights representing the contribution of each Eigenface image towards forming the original input image.

$$\omega_k = u_k^T(\tau - \psi) \quad (8)$$

The weights form the vector ω^T that represents the contribution of each Eigenface image towards forming the original input image.

$$\Omega^T = \{\omega_1, \omega_2, \dots, \omega_{M'}\} \quad (9)$$

For each training face image, we derive its corresponding vector of Eigenface weights, Ω_k , to be used for recognition. Turk and Pentland refer to this as a "face class".

$$\Omega_k = \{\omega_{k,1}, \omega_{k,2}, \dots, \omega_{k,M'}\} \quad (10)$$

$$\omega_{k,i} = u_i^T(\tau - \Psi) \quad (11)$$

Given the input image vector τ , we map it to the face class that provides the best description of the input image by calculating the Euclidean distance between its weight vector Ω and each of the face classes and select the face class yielding the minimum distance.

$$\epsilon_k^2 = ||(\Omega - \Omega_k)||^2 \quad (12)$$

Given this minimum distance and some preconfigured threshold Θ_k , we can determine whether the input image maps closely to one of the faces in our training set.

Even though we have identified which face in the training set most closely maps to the input image, many images will map to the training faces to some degree under Θ_k . The second part of classification, identification of faces, is described in the next section.

Identifying Faces

In order to ensure that the input image that we are classifying is an actual face (as opposed to Bart Simpson), Turk and Pentland propose a comparison between the original input image and a reconstructed projection image built from the vector of Eigenfaces u_k and their corresponding weights, ω_k . From this, the squared distance between the mean-adjusted input image and the reconstructed projection can determine how closely the input maps to our set of training faces (13).

$$\epsilon^2 = ||\Psi - \Psi_f||^2 \quad (13)$$

For some preconfigured threshold Θ , we can determine whether the input image projects onto "face space", allowing us to determine whether the image we are analyzing is actually a face. If the distance ϵ is above Θ , we should assume that the input image is not a face.

Implementation

The facial recognition engine is implemented in Scala 2.10.4 running on Java 1.7.0.79. Image processing logic uses common code pulled from the open source project Tinderbox (<https://github.com/crockpotveggies/tinderbox>). Eigenface generation and facial identification matrix computation uses the Apache Commons Math library.

Testing

The Faces94 collection maintained by the School of Computer Science and Electronic Engineering at the University of Essex was used as the pool of faces for both training and testing. For each individual subject, the Faces94 collection contained 16 images of that subject, each with minor variations in position, facial expression, etc. These images were included in the pool of faces. In addition, stock images of nonhuman faces (e.g. birds, clouds, cars, etc.) were included to test for face identification.

For each run, 30 human faces were randomly selected from the pool of faces to be used as the training set, and 30 human faces and nonface images were randomly selected to be used as the testing set. Given 153 individuals and some miscellaneous, most runs consisted of around 3-5 overlapping samples between the training set and testing set, as well as 2-3 images of nonhuman

faces.. This allows for more robust testing of both the facial recognition and identification algorithms, as each run will have a more diverse set of training/testing inputs and varying expected outputs.

Results

Based on empirical testing, the magnitude threshold for the minimum Θ_k was set to $2.0 * 10^{26}$ for projecting an input image onto a specific face class. Input images above this threshold are considered to be unrecognized. The maximum allowed distance Θ between an input image and its Eigenvalue projected representation was set to $1.0 * 10^{13}$. Input images with a distance above this threshold are determined to not be images of human faces.

Error rates were calculated over 30 runs, each of which tested against 30 randomly selected face/nonface images.

Of 900 test images, 884 were identified correctly, yielding an overall error rate of 1.78%. Of the 16 test images that were misidentified, 10 were nonface images identified as human face images and 6 were human face images identified as nonface images. Increasing the threshold Θ yielded fewer nonface images that were identified as human faces, but also resulted in more human face images being identified as nonhuman.

Of 837 human faces that were classified and identified as human faces, 655 were classified correctly as recognized or not recognized, yielding an overall error rate of 78.3%. Of the 182 human faces that were misclassified, 162 were faces which were part of the training set that were not recognized and 20 were faces which were not part of the training set that were mistakenly recognized. This larger proportion of false negatives is likely attributed to a combination of the variance between the training faces and the testing face for a given individual subject, as well as possibly the need for more than 10 Eigenface vectors to sufficiently recognize an individual. It is likely that the combination of increasing the number of Eigenfaces and increasing the magnitude threshold Θ_k would decrease the number of false negatives at the cost of an increased number of false positives.